

Machine Learning Assignment 3

Bohan YANG
Student ID: 2330016056

November 5, 2025

Question 1

Class +1 : (1, 2), (2, 3) Class -1 : (2, 1), (3, 2)

(1)

The separating hyperplane is

$$w^\top x + b = 0, \quad w \in \mathbb{R}^2, \quad b \in \mathbb{R}.$$

For the hard-margin SVM, the optimization problem is

$$\begin{aligned} & \min_{w,b} \quad \frac{1}{2} \|w\|_2^2 \\ \text{s.t.} \quad & y_i(w^\top x_i + b) \geq 1, \quad i = 1, \dots, 4. \end{aligned}$$

(2)

The Lagrangian is

$$\mathcal{L}(w, b, \alpha) = \frac{1}{2} \|w\|_2^2 - \sum_{i=1}^4 \alpha_i (y_i(w^\top x_i + b) - 1).$$

The KKT conditions:

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial w} = 0 & \Rightarrow w = \sum_{i=1}^4 \alpha_i y_i x_i, \\ \frac{\partial \mathcal{L}}{\partial b} = 0 & \Rightarrow \sum_{i=1}^4 \alpha_i y_i = 0. \end{aligned}$$

Substitute these into \mathcal{L} to obtain the dual optimization problem:

$$\begin{aligned}
\max_{\alpha \geq 0} \quad & \sum_{i=1}^4 \alpha_i - \frac{1}{2} \sum_{i=1}^4 \sum_{j=1}^4 \alpha_i \alpha_j y_i y_j x_i^\top x_j, \\
\text{s.t.} \quad & \sum_{i=1}^4 \alpha_i y_i = 0.
\end{aligned}$$

(3)

Assume all $\alpha_i = 0.5$ and that the point $(1, 2)$ is a support vector.

Then,

$$w = \sum_{i=1}^4 \alpha_i y_i x_i = 0.5 \left[(1, 2) + (2, 3) - (2, 1) - (3, 2) \right] = 0.5 (-2, 2) = (-1, 1).$$

For the support vector $(1, 2)$ with $y = +1$, the KKT equality holds:

$$y_i(w^\top x_i + b) = 1 \Rightarrow (-1, 1) \cdot (1, 2) + b = 1 \Rightarrow (-1 + 2) + b = 1 \Rightarrow b = 0.$$

Therefore, the optimal separating hyperplane is

$$w^\top x + b = 0 \Rightarrow (-1, 1) \cdot (x_1, x_2) + 0 = 0 \Rightarrow x_2 - x_1 = 0.$$

Question 2

(1)

Given training data

$$\{(x_i, y_i)\}_{i=1}^n, \quad x_i \in \mathbb{R}^d, \quad y_i \in \{+1, -1\}.$$

Let

$$w \in \mathbb{R}^d, \quad b \in \mathbb{R}, \quad \xi = [\xi_1, \dots, \xi_n]^\top \in \mathbb{R}^n, \quad \xi_i \geq 0.$$

The soft-margin SVM primal optimization problem is

$$\begin{aligned}
\min_{w, b, \xi} \quad & \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^n \xi_i \\
\text{s.t.} \quad & y_i(w^\top x_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n, \\
& \xi_i \geq 0, \quad i = 1, \dots, n.
\end{aligned}$$

With $C = 1$, this becomes

$$\min_{w, b, \xi} \frac{1}{2} \|w\|_2^2 + \sum_{i=1}^n \xi_i \quad \text{s.t. } y_i(w^\top x_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$

(2)

Introduce Lagrange multipliers

$$\begin{aligned}\alpha_i &\geq 0 \text{ for } y_i(w^\top x_i + b) \geq 1 - \xi_i \\ \mu_i &\geq 0 \text{ for } \xi_i \geq 0.\end{aligned}$$

The Lagrangian is

$$\mathcal{L}(w, b, \xi; \alpha, \mu) = \frac{1}{2} \|w\|_2^2 + \sum_{i=1}^n \xi_i - \sum_{i=1}^n \alpha_i (y_i(w^\top x_i + b) - 1 + \xi_i) - \sum_{i=1}^n \mu_i \xi_i.$$

Stationarity conditions:

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w} &= 0 \Rightarrow w = \sum_{i=1}^n \alpha_i y_i x_i, \\ \frac{\partial \mathcal{L}}{\partial b} &= 0 \Rightarrow \sum_{i=1}^n \alpha_i y_i = 0, \\ \frac{\partial \mathcal{L}}{\partial \xi_i} &= 0 \Rightarrow 1 - \alpha_i - \mu_i = 0 \Rightarrow \alpha_i + \mu_i = 1.\end{aligned}$$

Since $\mu_i \geq 0$, we have

$$0 \leq \alpha_i \leq 1 \quad (\text{generally } 0 \leq \alpha_i \leq C).$$

Substitute back:

$$\begin{aligned}\mathcal{L} &= \frac{1}{2} \|w\|_2^2 - \sum_{i=1}^n \alpha_i y_i w^\top x_i + \sum_{i=1}^n \alpha_i \\ &= -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^\top x_j + \sum_{i=1}^n \alpha_i,\end{aligned}$$

Dual problem:

$$\begin{aligned}\max_{\alpha} \quad & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^\top x_j, \\ \text{s.t.} \quad & \sum_{i=1}^n \alpha_i y_i = 0, \quad 0 \leq \alpha_i \leq 1.\end{aligned}$$

KKT conditions:

$$\begin{cases} w = \sum_i \alpha_i y_i x_i, \\ \sum_i \alpha_i y_i = 0, \\ 0 \leq \alpha_i \leq 1, \\ y_i(w^\top x_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0, \\ \alpha_i [y_i(w^\top x_i + b) - 1 + \xi_i] = 0, \\ (1 - \alpha_i) \xi_i = 0. \end{cases}$$

(3)

At optimality,

$$w^* = \sum_{i=1}^n \alpha_i^* y_i x_i.$$

For any support vector with $0 < \alpha_i^* < 1$,

$$y_i((w^*)^\top x_i + b^*) = 1.$$

Thus,

$$b^* = y_i - (w^*)^\top x_i.$$

The decision function is

$$f(x) = \text{sign}((w^*)^\top x + b^*), \quad \text{and the decision boundary is } (w^*)^\top x + b^* = 0.$$

(4)

In the Lagrangian, ξ_i appears only linearly:

$$\sum_i \xi_i - \sum_i \alpha_i \xi_i - \sum_i \mu_i \xi_i = \sum_i \xi_i (1 - \alpha_i - \mu_i).$$

Setting the derivative w.r.t. ξ_i to zero gives $1 - \alpha_i - \mu_i = 0$, or $\alpha_i + \mu_i = 1$.

Substituting this back eliminates all terms involving ξ_i , so ξ_i no longer appears in the dual function.

Its only effect is to constrain $\alpha_i \leq C (= 1)$, which forms the box constraint $0 \leq \alpha_i \leq C$.

Question 3

$$x_1 = (1, 2), \quad y_1 = +1; \quad x_2 = (2, 3), \quad y_2 = +1; \quad x_3 = (3, 1), \quad y_3 = -1; \quad x_4 = (4, 3), \quad y_4 = -1.$$

Polynomial kernel: $k(x_i, x_j) = (x_i^\top x_j + 1)^2$. Regularization $C = 1$.

(1)

$$\begin{aligned} x_1^\top x_1 &= 5, & x_1^\top x_2 &= 8, & x_1^\top x_3 &= 5, & x_1^\top x_4 &= 10, \\ x_2^\top x_2 &= 13, & x_2^\top x_3 &= 9, & x_2^\top x_4 &= 17, \\ x_3^\top x_3 &= 10, & x_3^\top x_4 &= 15, \\ x_4^\top x_4 &= 25. \end{aligned}$$

$$K = \begin{pmatrix} (5+1)^2 & (8+1)^2 & (5+1)^2 & (10+1)^2 \\ (8+1)^2 & (13+1)^2 & (9+1)^2 & (17+1)^2 \\ (5+1)^2 & (9+1)^2 & (10+1)^2 & (15+1)^2 \\ (10+1)^2 & (17+1)^2 & (15+1)^2 & (25+1)^2 \end{pmatrix} = \begin{pmatrix} 36 & 81 & 36 & 121 \\ 81 & 196 & 100 & 324 \\ 36 & 100 & 121 & 256 \\ 121 & 324 & 256 & 676 \end{pmatrix}.$$

(2)

Using the soft-margin SVM dual with kernel,

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^4} \quad & \sum_{i=1}^4 \alpha_i - \frac{1}{2} \sum_{i=1}^4 \sum_{j=1}^4 \alpha_i \alpha_j y_i y_j K_{ij}, \\ \text{s.t.} \quad & \sum_{i=1}^4 \alpha_i y_i = 0, \quad 0 \leq \alpha_i \leq C (= 1). \end{aligned}$$

At optimality, the decision function is

$$f(x) = \sum_{i=1}^4 \alpha_i^* y_i k(x_i, x) + b^*, \quad \hat{y}(x) = \text{sign}(f(x)).$$

(3)

Given

$$\alpha_1 = 0.0182, \quad \alpha_2 = 0.0068, \quad \alpha_3 = 0.0250, \quad \alpha_4 = 0,$$

and x_3 is a support vector with $y_3 = -1$.

For any margin support vector x_s with $0 < \alpha_s < C$,

$$y_s \left(\sum_{j=1}^4 \alpha_j y_j K_{js} + b \right) = 1.$$

Take $s = 3$:

$$-1 \left(\sum_{j=1}^4 \alpha_j y_j K_{j3} + b \right) = 1 \Rightarrow b = -1 - \sum_{j=1}^4 \alpha_j y_j K_{j3}.$$

Compute the sum using the 3rd column of K : $(36, 100, 121, 256)^\top$.

$$\sum_{j=1}^4 \alpha_j y_j K_{j3} = 0.0182 \cdot (+1) \cdot 36 + 0.0068 \cdot (+1) \cdot 100 + 0.0250 \cdot (-1) \cdot 121 + 0 \cdot (-1) \cdot 256 = -1.6898.$$

Hence

$$b = -1 - (-1.6898) = 0.6898.$$

(4)

Evaluate

$$f(x_5) = \sum_{i=1}^4 \alpha_i y_i k(x_i, x_5) + b.$$

First compute $k(x_i, x_5) = (x_i^\top x_5 + 1)^2$:

$$\begin{aligned} x_1^\top x_5 &= 1 \cdot 2 + 2 \cdot 1 = 4 \Rightarrow k(x_1, x_5) = (4 + 1)^2 = 25, \\ x_2^\top x_5 &= 2 \cdot 2 + 3 \cdot 1 = 7 \Rightarrow k(x_2, x_5) = (7 + 1)^2 = 64, \\ x_3^\top x_5 &= 3 \cdot 2 + 1 \cdot 1 = 7 \Rightarrow k(x_3, x_5) = (7 + 1)^2 = 64, \\ x_4^\top x_5 &= 4 \cdot 2 + 3 \cdot 1 = 11 \Rightarrow k(x_4, x_5) = (11 + 1)^2 = 144. \end{aligned}$$

Thus

$$\begin{aligned}f(x_5) &= 0.0182 \cdot (+1) \cdot 25 + 0.0068 \cdot (+1) \cdot 64 + 0.0250 \cdot (-1) \cdot 64 + 0 \cdot (-1) \cdot 144 + 0.6898 \\&= 0.455 + 0.4352 - 1.6 + 0 + 0.6898 \\&= -0.0200 \text{ (approximately).}\end{aligned}$$

Therefore,

$$\boxed{\hat{y}(x_5) = \text{sign}(f(x_5)) = -1.}$$

DS4023_assignment3

November 5, 2025

1 DS4023 Assignment 3: SVM and Ensemble Learning(45 pts)

1.0.1 Q1: Soft-margin SVM (9 pts)

You are given a dataset (svm_soft.mat) containing 2D points ['X'] from two classes with corresponding label (0 or 1) in column ['y']. Please visualize the datapoints and classify them using linear SVM with soft margin. You may assume that the penalty term C is set to 1 in $\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_i \epsilon_i$.

```
[1]: """
Part 1: Visualize the datapoints. Your plot should look similar to the sample
       ↵one.

"""

import pandas as pd
import matplotlib.pyplot as plt
from scipy.io import loadmat

def plot_svm_soft(boundary=False, svc=None):
    raw_data = loadmat('data/svm_soft.mat')

    data = pd.DataFrame(raw_data['X'], columns=['X1', 'X2'])
    data['y'] = raw_data['y']

    pos = data['y'] == 1
    neg = data['y'] == 0

    plt.figure(figsize=(8, 6))
    plt.scatter(data.loc[pos, 'X1'], data.loc[pos, 'X2'],
                marker='x', color='red', label='Positive')
    plt.scatter(data.loc[neg, 'X1'], data.loc[neg, 'X2'],
                marker='o', color='blue', label='Negative')

    if boundary:
        X = data[['X1', 'X2']].values
        y = data['y'].values

        x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
```

```

y_min, y_max = X[:, 1].min()-0.5, X[:, 1].max()+0.5
xx, yy = np.meshgrid(
    np.linspace(x_min, x_max, 300),
    np.linspace(y_min, y_max, 300)
)


Z = svc.decision_function(
    np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)


plt.contour(xx, yy, Z, levels=[0],
            linestyles=['-'], colors='lightblue')

plt.legend()
plt.show()

```

plot_svm_soft()





[2]:

```
"""
Part 2: Classify the datapoints and visualize the decision boundary.
You should print out the coefficients and the interception for the
linear classifier.
Your plot should look like the sample one.
"""

from sklearn import svm
import numpy as np
```

```
def classify_svm_soft():
    raw_data = loadmat('data/svm_soft.mat')
    data = pd.DataFrame(raw_data['X'], columns=['X1', 'X2'])
    data['y'] = raw_data['y']

    svc = svm.LinearSVC(C=1, loss='hinge', max_iter=1_000_000)
    svc.fit(data[['X1', 'X2']], data['y'])

    print('Coefficients (w):', svc.coef_[0])
    print('Intercept (b):', svc.intercept_[0])


    plot_svm_soft(boundary=True, svc=svc)

    return svc


classify_svm_soft()
```

```
Coefficients (w): [0.59153686  0.81825054]
Intercept (b): -4.11954105475478
/opt/homebrew/anaconda3/lib/python3.11/site-
```

```
packages/sklearn/utils/validation.py:2749: UserWarning: X does not have valid  
feature names, but LinearSVC was fitted with feature names  
warnings.warn(
```



```
[2]: LinearSVC(C=1, loss='hinge', max_iter=1000000)
```



1.0.2 Q2: Kernel SVM (9 pts)

You are given a dataset (svm_kernel.mat) containing 2D points ['X'] from two classes with corresponding label (0 or 1) in column ['y']. Please visualize the datapoints and classify them using linear SVM with Gaussian kernel.

```
[8]: """
Part 1: Visualize the datapoints. Your plot should look similar to the sample
↓one.

"""

import pandas as pd
import matplotlib.pyplot as plt
from scipy.io import loadmat
import seaborn as sb



def plot_svm_kernel():
    global data_svm_kernel

    data_svm_kernel = loadmat('data/svm_kernel.mat')
    X2 = data_svm_kernel['X']
    y2 = data_svm_kernel['y'].ravel()

    data = pd.DataFrame(data_svm_kernel['X'], columns=['X1', 'X2'])
    data['y'] = data_svm_kernel['y']

    plt.figure(figsize=(12, 6))
    sb.scatterplot(
        data=data,
        x='X1', y='X2',
        hue='y',
        style='y',
        palette={1: '#1f77b4', 0: '#ff7f0e'},
        markers={1: 'X', 0: 'o'},
        s=50,
    )
    plt.xlabel('')
    plt.ylabel('')
    plt.legend(labels=['Positive', 'Negative'])
    plt.tight_layout()
    plt.show()

plot_svm_kernel()
```



[6] :

```
"""
Part 2. Classify the data points using Gaussian kernel.
Visualize the probabilities of the classification results.
If your classification is correct, your plot should look similar to the
sample one.
"""
```

```
def classify_svm_kernel():
    global data_svm_kernel
    data = loadmat('data/svm_kernel.mat')
    X = data['X']
```


```


y = data['y'].flatten()
data_svm_kernel = {
    'X1': X[:, 0],
    'X2': X[:, 1]
}
svc = svm.SVC(kernel='rbf', C=1000, gamma=100, probability=True)
svc.fit(X, y)
probabilities = svc.predict_proba(X)[:, 0]
return probabilities

probabilities = classify_svm_kernel()

fig, ax = plt.subplots(figsize=(12, 8))
ax.scatter(data_svm_kernel['X1'], data_svm_kernel['X2'],
           s=30, c=probabilities, cmap='Reds')
plt.show()

```





1.0.3 Q3: SVM in Email Spam Detection (9 pts)

Many email services today provide spam filters that are able to classify emails into spam and non-spam email with high accuracy. Please write an SVM classifier to detect spam emails. There are two data files that you need to load, one for training, i.e., “email_train.mat”, and the other for testing, i.e., “email_test.mat”. Your model should train on the training dataset and be applied to the test dataset to obtain the spam detection accuracy. **Your classifier should have an accuracy higher than 90% in both training and testing. Otherwise, every 10% of accuracy decrease leads to a 2 pts deduction in marks.** For example, an accuracy of [80%, 90%) will get deducted for 2 pts overall, and [70%, 80%) will get deducted for 4 pts.

```
[ ]: from scipy.io import loadmat

def spam_detector_svm():
    train_data = loadmat('data/email_train.mat')
    X = train_data['X']
    y = train_data['y'].flatten()

    test_data = loadmat('data/email_test.mat')
    Xtest = test_data['Xtest']
    ytest = test_data['ytest'].flatten()

    svc = svm.SVC(kernel='rbf', C=1000, gamma='scale')
    svc.fit(X, y)
```

```

# Do not remove.
print('Training accuracy = {:.0}%'.format(
    np.round(svc.score(X, y) * 100, 2)))
print('Test accuracy = {:.0}%'.format(
    np.round(svc.score(Xtest, ytest) * 100, 2)))

spam_detector_svm()

```

Training accuracy = 100.0%
Test accuracy = 98.0%

1.0.4 Q4: Bagging (9 pts)

You are given a tumor dataset `tumor.csv`. The dataset contains instances with 9 features. Each instance is labeled in either benign or malignant classes (0 for benign, 1 for malignant in last column). The dataset only contains numeric values and has been normalized. Use `RandomForestClassifier` to calculate the mean accuracy by using 8-folds cross-validation. Number of trees set to 50.

```

[ ]: import numpy as np
import pandas as pd
from sklearn import model_selection
from sklearn.ensemble import RandomForestClassifier


def classify_bagging():
    rand_seed = 50 # Do not change

    df = pd.read_csv('data/tumor.csv')

    X = df.iloc[:, :-1].values
    y = df.iloc[:, -1].values

    rf = RandomForestClassifier(n_estimators=50, random_state=rand_seed)
    results = model_selection.cross_val_score(rf, X, y, cv=8)

    # Do not remove.
    print("Performance of random forest:{:.2f} % ({:.2f})".format(results.mean(), results.std()))

classify_bagging()

```

Performance of random forest:0.965713

1.0.5 Q5: AdaBoost (9 pts)

Use `AdaBoostClassifier` to classify the same dataset `tumor.csv`. Compute the mean accuracy by using 8-folds cross-validation and 50 estimators.

```
[ ]: from sklearn.ensemble import AdaBoostClassifier

def classify_adaboost():
    rand_seed = 50 # Do not change

    df = pd.read_csv('data/tumor.csv')
    X = df.iloc[:, :-1].values
    y = df.iloc[:, -1].values

    ada = AdaBoostClassifier(n_estimators=50, random_state=rand_seed)
    cv = StratifiedKFold(n_splits=8, shuffle=True, random_state=rand_seed)
    results = cross_val_score(
        ada, X, y, cv=cv, scoring='accuracy', n_jobs=-1)

    # Do not remove.
    print("Performance of adaboost:%f" % (results.mean()))

classify_adaboost()
```

```
Performance of adaboost:0.954219
```

```
[ ]:
```